

BH Evidence Lens

From assumption to data-backed decision

A HUMAN-LED AI AGENT

BH Evidence Lens builds and tests project assumptions using approved BH knowledge and current project evidence, then explores how future technologies, behavioural change and alternative scenarios could affect those assumptions. The agent strengthens the evidence; the consultant decides.

PROBLEM

Evidence is fragmented and difficult to challenge.

SOLUTION

A repeatable, traceable evidence workflow.

OUTCOME

More defensible decisions with consultant validation.

1. Executive Summary

BH Evidence Lens is a human-led AI agent designed to make project assumptions visible, testable and challengeable. It brings together approved BH knowledge, project information, site evidence and external research, then analyses how future technologies, behavioural change, policy shifts and alternative scenarios could affect the project assumptions through a repeatable workflow. The agent supports evidence gathering and analysis, while the consultant remains responsible for interpretation, validation and the final decision.

The opportunity is not limited to faster searching or drafting. The strategic value is a consistent evidence trail that shows what information was used, why it is comparable, where uncertainty remains and what must be true for a proposed target or scenario to succeed.

The challenge

- Project assumptions are informed by multiple sources, including personal experience, past reports, client targets, external research, site evidence, programme and phasing.
- The underlying evidence is often fragmented, hidden in different files or reconstructed differently by different consultants.
- When an assumption cannot be traced, it cannot be properly challenged, repeated or confidently explained to a client.

The proposed response

Evidence inputs	Configured workflow	Consultant output
• Approved BH knowledge • Project data • Site evidence • External research	• Retrieve and filter • Challenge assumptions • Test future scenarios • Assess technology impacts • Record evidence gaps	• Source trail • Ranges and sensitivities • Validation gate • Defensible recommendation

Why it is different

The differentiator is the workflow, not only the base AI model. BH Evidence Lens combines approved sources, access controls, reusable retrieval, project taxonomy, governed external research, scenario logic and consultant validation gates. This turns general AI capability into a repeatable BH decision-support method.

Intended users and application

The initial demonstration is transport-focused, but the underlying approach is applicable wherever project decisions depend on assumptions and precedent. The agent can also explore how emerging technologies, behavioural change, policy shifts and alternative future conditions may alter demand, capacity or delivery requirements across transport, urban strategy, infrastructure, climate and resilience, and buildings and places.

Future assumptions, technology and scenario analysis

The agent will not treat today's evidence as a fixed forecast. It will use the current evidence base as the starting point, identify the assumptions that may change over time and test a range of plausible future conditions. The purpose is to show how a recommendation performs under different futures, not to claim that one future can be predicted with certainty.

- Future assumptions: make expected changes in population, behaviour, land use, policy, phasing and demand explicit.
- Technology pathways: assess how new mobility, automation, digital services, energy systems or other relevant technologies could change the original assumption.
- Alternative scenarios: compare baseline, accelerated-change and constrained-delivery conditions using consistent drivers.
- Sensitivity and triggers: identify which variables have the greatest influence and what evidence should trigger a future review.
- Decision support: present robust actions, flexible actions and assumptions requiring consultant or client approval.

The output remains evidence-led. Each future scenario should state its source, date, assumption, uncertainty and consultant validation requirement.

2. How BH Evidence Lens Works

01	RETRIEVE	Find relevant BH precedent, project evidence and approved external sources.
02	FILTER	Use a project taxonomy to select genuinely comparable evidence.
03	CHALLENGE	Expose the initial assumption, confidence level and missing information.
04	TEST	Apply consistent drivers, future technology assumptions, alternative scenarios and sensitivity ranges.
05	TRACE	Record source, date, assumption and uncertainty throughout the process.
06	VALIDATE	Place the consultant at the final review and approval gate.

3. Demonstrated Application: NoRA Mobility Strategy

The concept was applied to the NoRA Mobility Strategy, where the proposed mobility outcomes included 40%+ public transport, 50%+ sustainable internal trips and 70%+ trips internalised. The purpose was not to invent new targets, but to reconstruct the evidence path and test the conditions required to support those targets.

40%+	50%+	70%+
PUBLIC TRANSPORT	SUSTAINABLE INTERNAL TRIPS	TRIPS INTERNALISED

What the evidence workflow demonstrated

- A benchmark base of 150 cities could be filtered to 11 relevant comparators rather than treated as one undifferentiated dataset.
- External public transport share and internal sustainable travel needed to be separated because they represent different project outcomes.
- The targets depended on enabling conditions such as transit supply, density, access, parking and phasing.
- The output could distinguish what was approved, what had been modelled and what evidence was still missing.
- The final package could present a range, sensitivity, source trail and review gate instead of a single unsupported “magic number”.

4. Business Value and Pilot Logic

The value proposition is to reuse the evidence rather than repeatedly recreate the effort. Current work often involves finding the correct files, extracting benchmarks, resolving conflicting assumptions, rebuilding charts and creating a source trail near the end. BH Evidence Lens is intended to replace these repeated steps with a reusable retrieval layer, comparator taxonomy, scenario template and structured review gate.

Illustrative value models

Model	Inputs	Calculated potential	Evidence status
Complex evidence questions	100 questions/year; 6 hours saved/question	600 consulting hours/year	Illustrative pilot assumptions
Wider employee adoption	100 employees; 3 questions/day; 22 days/month; 10 minutes saved/question	1,100 hours/month; 13,200 hours/year	Illustrative adoption assumptions
Shared agent cost comparison	6,600 questions/month; PAYG credit assumptions	Deck scenario indicates lower cost than blanket licensing	Requires verification of BH pricing and agent design

Important interpretation

The time and cost figures are forecast scenarios, not measured savings. The pilot must replace these assumptions with observed completion times, correction effort, user adoption and actual Copilot Studio credit consumption.

Recommended pilot

Run a controlled pilot using a defined set of approved BH documents and one or two repeatable evidence questions. Compare conventional delivery with agent-assisted delivery, while recording the same quality, time, traceability and cost measures for both approaches.

Pilot success measures

- Retrieval accuracy and relevance of returned sources
- Percentage of material statements linked to an identifiable source
- Time taken before and after using the agent
- Consultant review and correction effort
- Consistency when the same question is repeated
- Unsupported or outdated claims detected
- Actual credits consumed per simple and complex request
- Active users, repeat usage and user-rated usefulness
- Quality of future scenarios, including evidence, transparency of assumptions and usefulness of sensitivities

5. Evidence Package Summary

The evidence package should distinguish demonstrated project evidence from proposed configuration and unvalidated business assumptions. This separation protects the credibility of the submission and gives the judging panel a clear basis for assessing readiness.

Evidence category	What is available	Current status
Problem and concept	Defined challenge, human-led principle and cross-disciplinary use case	Demonstrated in pitch deck
Workflow design	Retrieval, filtering, assumption challenge, scenario testing, traceability and validation	Configured concept
NoRA case evidence	Project KPIs, 150-city benchmark base, 11 comparators and testing framework	Supported by named internal project sources
Prototype approach	Copilot Studio agent using selected knowledge, instructions, skills and governance	Proposed implementation
Business value	Time-release and licence-cost scenarios	Illustrative; pilot validation required
Risk and governance	Source permissions, traceability, consultant approval and evidence maturity	Core design requirement
Future scenario analysis	Future assumptions, technology pathways, alternative scenarios, sensitivities and review triggers	Proposed analytical capability; pilot validation required

Source register

20260724_NoRA Urban framework July BoD section_v5.pptx: Project targets: 40%+ public transport, 50%+ sustainable internal trips and 70%+ internalised trips.

NoRA_Mode_Share_Past_Work_Recap.pptx: Benchmark evidence base: 150 cities, 11 shortlisted comparators and three comparator groups.

NoRA_Mode_Share_Benchmark_Deck.pptx: Evidence-testing framework: segment, test, localise, enable and monitor.

BH AI Infra Strategy, 17 September 2026: Internal need for a shared Teams agent to reduce searching for and recreating existing work.

Training: A BH intro to AI, Session 01, 9 September 2026: Internal example indicating that contextual search can materially reduce retrieval time.

Microsoft Copilot Studio pricing and billing guidance: Public basis used in the pitch deck for preliminary credit and operating-cost scenarios.

Conclusion

BH Evidence Lens addresses a practical consulting challenge: important assumptions are often supported by fragmented evidence that cannot be easily repeated or challenged. The concept shows how a configured AI agent could convert BH knowledge and project information into a transparent evidence workflow, while also testing how future technologies and alternative scenarios may change current assumptions, without replacing professional judgement. The NoRA case demonstrates the relevance of the approach; the next step is a measured pilot to validate accuracy, time release, cost and scalability.